


Data Mining and Data Warehousing

2. Machine Learning I: Supervised Learning



Universitatea
Transilvania
din Braşov

FACULTATEA DE INGINERIE ELECTRICĂ
ŞI ŞTIINŢA CALCULATOARELOR

Drd. Horia Modran

Contact: horia.modran@unitbv.ro / modranhoria@gmail.com

Tel: 0770171577

2022 - 2023



What is learning ?

- Herbert Simon: "Learning is any process by which a system improves performance from experience."
- "A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E ."



– Tom Mitchell




Machine Learning

- Machine learning: study of computer algorithms that can improve automatically through experience and by use of data
- how to acquire a model on the basis of data / experience
 - learning parameters (e. g. probabilities)
 - learning structure (e. g. BN graphs)
 - learning hidden concepts (e. g. clustering)



Machine Learning and Data Mining

- Machine learning and data mining often employ the same methods and overlap significantly
- Machine learning focuses on prediction, based on known properties learned from the training data
- Data mining focuses on the discovery of (previously) unknown properties in the data (analysis step of knowledge discovery in databases)
- Data mining uses many machine learning methods, but with different goals





Machine Learning Areas

- Supervised Learning: data and corresponding labels are given
- Unsupervised Learning: only data is given, no labels provided
- Semi-supervised Learning: some (if not all) labels are present
- Reinforcement Learning: an agent interacting with the world makes observations, takes actions, and is rewarded or punished; it should learn to choose actions in such a way as to obtain a lot of reward



Supervised learning

- Supervised learning (SL) is the machine learning task of learning a function that maps an input to an output based on example input-output pairs
- it infers a function from labeled training data
- each example consists of an input object and a desired output
- an algorithm analyzes the training data and produces an inferred function, used for mapping new examples
- optimal scenario: correctly determine the class labels for unseen instances
- statistical quality - measured through generalization error



Important Concepts

- **Data:** labeled instances $\langle x, y \rangle$, e. g. emails marked spam/not spam -> split into training/(hold-out)/test set
- **Features:** attribute-value pairs which characterize each x
- **Experimentation cycle**
 - learn parameters (e. g. model probabilities) on training set
 - (Tune hyper-parameters on held-out set)
 - compute accuracy of test set
- **Evaluation**
 - accuracy: fraction of instances predicted correctly
- **Overfitting and generalization**
 - want a classifier which does well on test data



Example: Spam filter

Input: email

Output: spam/ham

Setup:

- Get a large collection of example emails, each labeled "spam" or "ham"
- Note: someone has to hand label all this data!
- Want to learn to predict labels of new, future emails

Features: The attributes used to make the ham / spam decision

- Words: FREE!
- Text Patterns: \$dd, CAPS
- Non-text: SenderInContacts
- ...



Dear Sir.

First, I must solicit your confidence in this transaction, this is by virtue of its nature as being utterly confidential and top secret. ...



TO BE REMOVED FROM FUTURE MAILINGS, SIMPLY REPLY TO THIS MESSAGE AND PUT "REMOVE" IN THE SUBJECT.

99 MILLION EMAIL ADDRESSES FOR ONLY \$99



Ok, I know this is blatantly OT but I'm beginning to go insane. Had an old Dell Dimension XPS sitting in the corner and decided to put it to use, I know it was working pre being stuck in the corner, but when I plugged it in, hit the power nothing happened.



Example: Digit Recognition

Input: images / pixel grids

Output: a digit 0-9

Setup:

- Get a large collection of example images, each labeled with a digit
- Note: someone has to hand label all this data!
- Want to learn to predict labels of new, future digit images

Features: The attributes used to make the digit decision

- Pixels: (6,8)=ON
- Shape Patterns: NumComponents, AspectRatio, NumLoops
- ...



0



1



2



1



??



Classification Examples

- In classification, we predict labels y (classes) for inputs x
- examples:
 - OCR (input: images, classes: characters)
 - Medical diagnosis (input: symptoms, classes: diseases)
 - Automatic essay grader (input: document, classes: grades)
 - Fraud detection (input: account activity, classes: fraud / no fraud)
 - Customer service email routing
 - recommended articles in a newspaper, books
 - DNA and protein sequence identification
 - financial investments





Inductive learning

- Simplest form: learn a function from examples
 - f is the target function \rightarrow an example is a pair $(x, f(x))$
- pure induction task:
 - given a collection of examples of f , return a function h that approximates f .
 - find a hypothesis h , such that $h \approx f$, given a training set of examples
- this is a highly simplified model of real learning:
 - ignores prior knowledge
 - assumes examples are given



Inductive learning

- construct/adjust h to agree with f on training set
 - h is consistent if it agrees with f on all examples
 - e. g. curve fitting




Ockham's razor: prefer the simplest hypothesis consistent with data



Generalization


- Hypotheses must generalize to correctly classify instances not in the training data.
- simply memorizing training examples is a consistent hypothesis that does not generalize.
- Occam's razor:
 - finding a simple hypothesis helps ensure generalization





Training error vs test error

- Low bias/high variance – **overfitting** ("fitting" the training set)
- High bias/low variance – **underfitting** (the model cannot capture the structure of the data)






Classification/Regression

- Learning a discrete function: **classification**
 - boolean classification:
 - each example is classified as true(positive) or false(negative)
 - can also predict multiple classes (3, 4, 5...)
- Learning a continuous function: **Regression**
 - E.g. Linear Regression, Logistic regression



Classification

- Model construction: describing a set of predetermined classes
 - each tuple/sample is assumed to belong to a predefined class, as determined by the class label
 - the set of tuples used for model construction is training set
 - the model is represented as classification rules, decision trees, or mathematical formulae
- Model usage: for classifying future or unknown objects
 - estimate accuracy of the model
 - if the accuracy is acceptable, use the model to classify data tuples whose class labels are not known



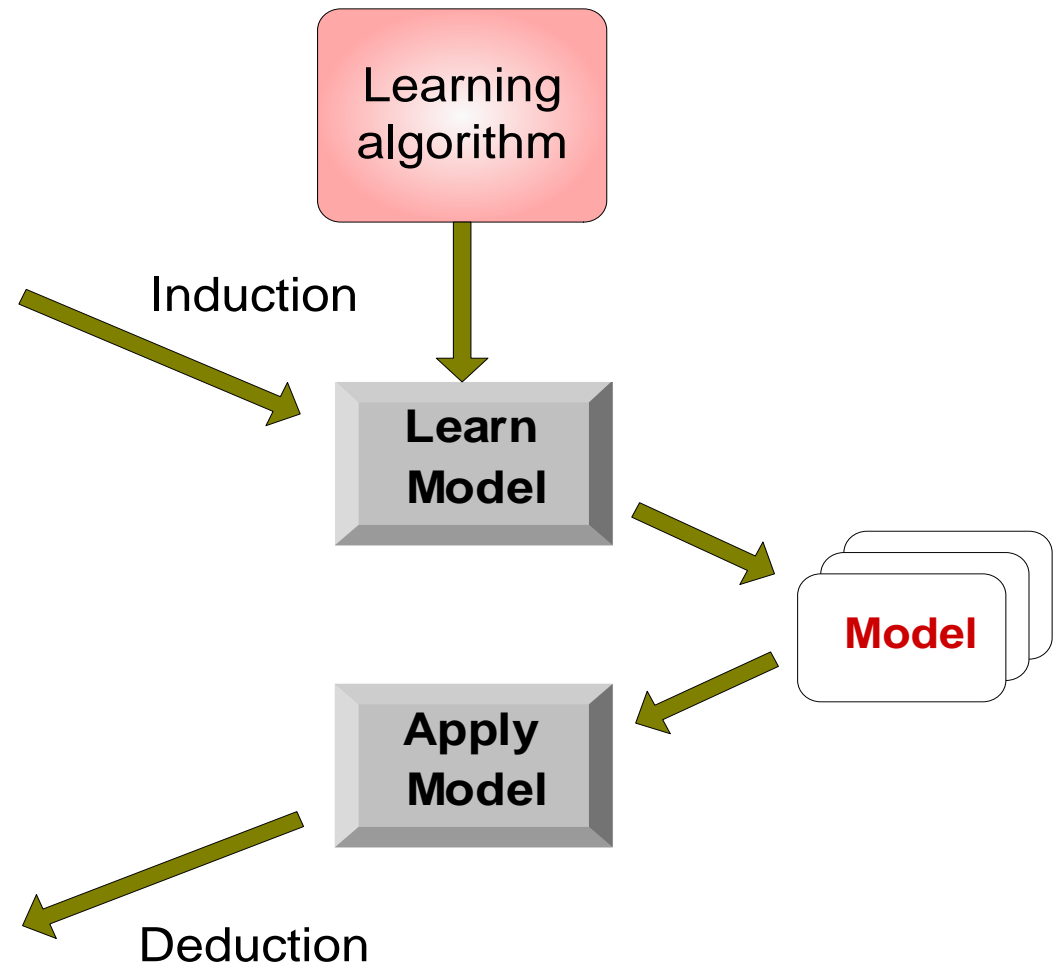
Classification example

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set





Data preparation

■ Data cleaning

- preprocess data in order to reduce noise and handle missing values

■ Relevance analysis (feature selection)

- remove the irrelevant or redundant attributes

■ Data transformation

- generalize data to (higher concepts, discretization)
- normalize attribute values






Classification techniques

- Decision Tree based Methods
 - Random Forests
- Rule-based Methods
- Naive Bayes
- Bayesian Belief Networks
- Support Vector Machines
- and many more...



Decision trees

- Example Problem: decide whether to wait for a table at a restaurant, based on the following attributes:
 - alternate: is there an alternative restaurant nearby?
 - Bar: is there a comfortable bar area to wait in?
 - Fri/Sat: is today Friday or Saturday?
 - hungry: are we hungry?
 - patrons: number of people in the restaurant (None, Some, Full)
 - price: price range (\$, \$\$, \$\$\$)
 - raining: is it raining outside?
 - reservation: have we made a reservation?
 - type: kind of restaurant (French, Italian, Thai, Burger)
 - waitEstimate: estimated waiting time (0-10, 10-30, 30-60, >60)



Feature base representation

- examples described by feature(attribute) values (Boolean, discrete, continuous)
- e.g., situations where I will/won't wait for a table:

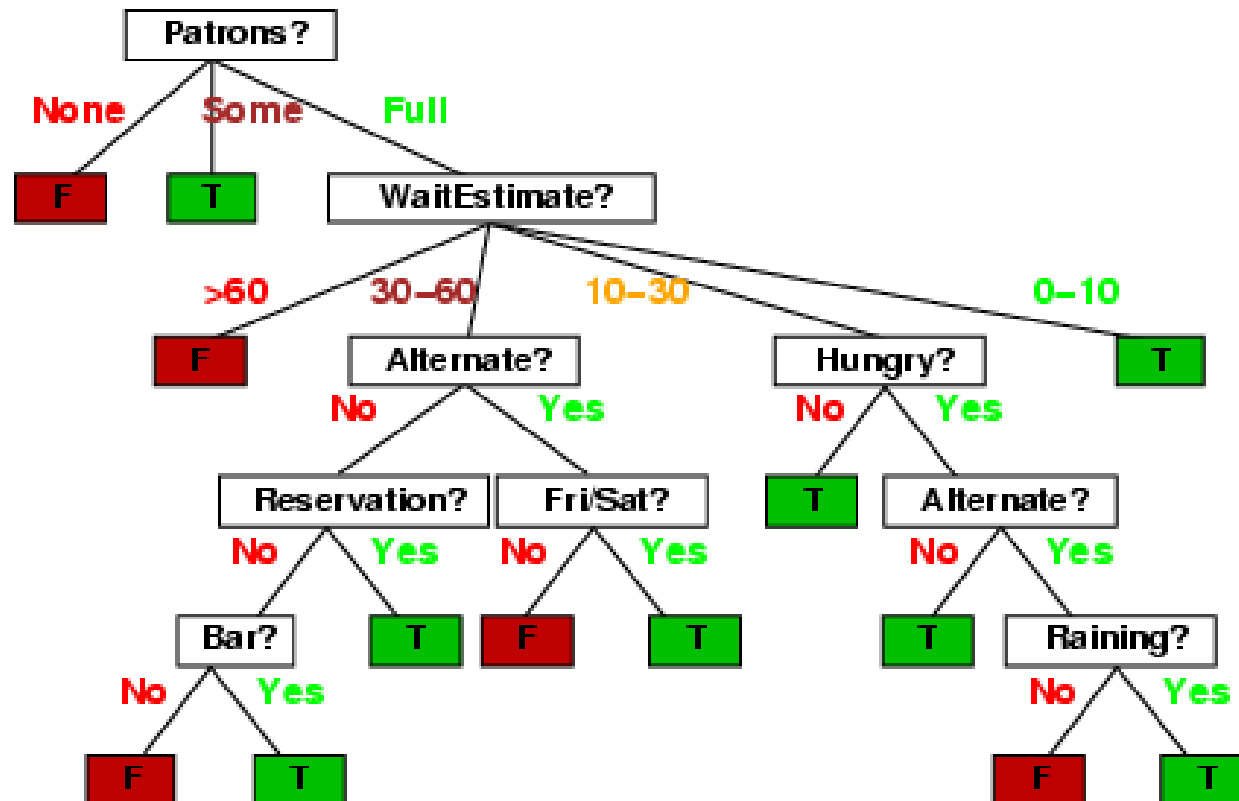
Example	Attributes										Target
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>Wait</i>
X_1	T	F	F	T	Some	\$\$\$	F	T	French	0–10	T
X_2	T	F	F	T	Full	\$	F	F	Thai	30–60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0–10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10–30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$\$	T	T	Italian	0–10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0–10	F
X_8	F	F	F	T	Some	\$\$	T	T	Thai	0–10	T
X_9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10–30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0–10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30–60	T

Classification of examples is positive (T) or negative (F)



Decision trees

- one possible representation for hypotheses
- e.g., here is the “true” tree for deciding whether to wait:






Expressivness

- Decision trees can express any function of the input attributes.
- e.g., for Boolean functions, truth table row \rightarrow path to leaf:

A	B	A xor B
F	F	F
F	T	T
T	F	T
T	T	F



- trivially, there is a consistent decision tree for any training set with one path to leaf for each example (unless f nondeterministic in x) but it probably won't generalize to new examples



Decision tree algorithm

- Principle – greedy algorithm
 - tree is constructed in a top-down recursive divide-and-conquer manner
- Iterations - at start, all the training tuples are at the root
 - tuples are partitioned recursively
 - test attributes are selected on the basis of a heuristic or statistical measure (e. g., information gain)
- Stopping conditions
 - all samples for a given node belong to the same class
 - there are no remaining attributes for further partitioning
 - there are no samples left



Advantages/disadvantages

■ Advantages

- **easy** to construct/implement
- **extremely fast** at classifying unknown records
- **accuracy** is comparable to other classification techniques for many simple data sets

■ Disadvantages

- computationally expensive to train
- some decision trees can be overly complex that do not generalize the data well
- less expressivity




K-Nearest Neighbor (KNN)

- simple, but a very powerful classification algorithm
- classifies based on a similarity measure
- non-parametric
- lazy learning
 - does not “learn” until the test example is given
 - whenever we have a new data to classify, we find its K nearest neighbors from the training data



KNN: Classification

- classified by “MAJORITY VOTES” for its neighbor classes
- assigned to the most common class amongst its K nearest neighbors (by measuring “distant” between data)





KNN: Steps


Step 1: Determine parameter K = number of nearest neighbors

Step 2: Calculate the distance between the query-instance and all the training examples.

Step 3: Sort the distance and determine nearest neighbors based on the k -th minimum distance.

Step 4: Gather the category Y of the nearest neighbors.

Step 5: Use simple majority of the category of nearest neighbors as the prediction value of the query instance.





KNN: Example




Table 1. Euclidean distance matrix D listing all possible pairwise Euclidean distances between 19 samples.

	x ₁	x ₂	x ₃	x ₄	x ₅	x ₆	x ₇	x ₈	x ₉	x ₁₀	x ₁₁	x ₁₂	x ₁₃	x ₁₄	x ₁₅	x ₁₆	x ₁₇	x ₁₈	x ₁₉
x ₂	1.5																		
x ₃	1.4	1.6																	
x ₄	1.6	1.4	1.3																
x ₅	1.7	1.4	1.5	1.5															
x ₆	1.3	1.4	1.4	1.5	1.4														
x ₇	1.6	1.3	1.4	1.4	1.5	1.8													
x ₈	1.5	1.4	1.6	1.3	1.7	1.6	1.4												
x ₉	1.4	1.3	1.4	1.5	1.2	1.4	1.3	1.5											
x ₁₀	2.3	2.4	2.5	2.3	2.6	2.7	2.8	2.7	3.1										
x ₁₁	2.9	2.8	2.9	3.0	2.9	3.1	2.9	3.1	3.0	1.5									
x ₁₂	3.2	3.3	3.2	3.1	3.3	3.4	3.3	3.4	3.5	3.3	1.6								
x ₁₃	3.3	3.4	3.2	3.2	3.3	3.4	3.2	3.3	3.5	3.6	1.4	1.7							
x ₁₄	3.4	3.2	3.5	3.4	3.7	3.5	3.6	3.3	3.5	3.6	1.5	1.8	0.5						
x ₁₅	4.2	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	1.7	1.6	0.3	0.5					
x ₁₆	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	1.6	1.5	0.4	0.5	0.4				
x ₁₇	5.9	6.2	6.2	5.8	6.1	6.0	6.1	5.9	5.8	6.0	2.3	2.3	2.5	2.3	2.4	2.5			
x ₁₈	6.1	6.3	6.2	5.8	6.1	6.0	6.1	5.9	5.8	6.0	3.1	2.7	2.6	2.3	2.5	2.6	3.0		
x ₁₉	6.0	6.1	6.2	5.8	6.1	6.0	6.1	5.9	5.8	6.0	3.0	2.9	2.7	2.4	2.5	2.8	3.1	0.4	

$$d(x, y) = \sqrt{\sum_{i=1}^n (y_i - x_i)^2}$$

n– number of dimensions (in our case – 2)



Pros and cons

■ Pros

- learning and implementation is extremely simple and intuitive
- flexible decision boundaries

■ Cons

- irrelevant or correlated features have high impact and must be eliminated
- difficult to handle high dimensionality
- computational costs: memory and classification time computation




Correlation

- linear association between two variables
- show how to determine both the nature and strength of relationship between two variables
- correlation lies between -1 to $+1$
- zero correlation indicates that there is no relationship between the variables
- Pearson correlation coefficient
 - most familiar measure of dependence between two quantities




Linear Regression



Samples with ONE independent variable



Samples with TWO independent variables



Given examples $(x_i, y_i)_{i=1 \dots n}$
Predict y_{n+1} given a new point x_{n+1}





Linear Regression

- how to represent the data as a vector/matrix
- we assume a model

$$\mathbf{y} = b_0 + \mathbf{bX} + \epsilon$$

, where b_0 and \mathbf{b} are intercept and slope, known as coefficients or parameters. ϵ is the error term (typically assumes that $\epsilon \sim N(\mu, \sigma^2)$)

- Simple Linear Regression
 - a single independent variable is used to predict
- Multiple linear regression
 - two or more independent variables are used to predict



Linear Regression

- find the optimal coefficient vector \mathbf{b} that makes the most similar observation: $\mathbf{y} = \mathbf{X}\mathbf{b} + \mathbf{e}$ (vector multiplication)

$$\begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_{11} & \cdots & x_{1p} \\ 1 & \vdots & \ddots & \vdots \\ 1 & x_{n1} & \cdots & x_{np} \end{bmatrix} \begin{bmatrix} b_0 \\ \vdots \\ b_p \end{bmatrix} + \begin{bmatrix} e_1 \\ \vdots \\ e_n \end{bmatrix}$$

- need to minimize the error (sum squared error)

$$\min J(\mathbf{b}) = \sum_{i=1}^n (y_i - \mathbf{x}_{i,*} \mathbf{b})^2$$



Linear regression with categorical variables

- we assumed that all variables are continuous variables
- categorical variables:
 - ordinal variables - encode data with continuous values
 - Evaluation: Excellent (5), Very good (4), Good (3), Poor (2), Very poor (1)
 - nominal variables - use dummy variables
 - department: Computer, Biology, Physics

	Computer	Biology	Physics
Computer	1	0	0
Biology	0	1	0
Physics	0	0	1



Linear regression for classification

- for binary classification
 - encode class labels as $y=0,1$ or $\{-1,1\}$
 - apply formula: $\mathbf{y} = \mathbf{X} * \mathbf{b} + \mathbf{e}$
 - check which class the prediction is closer to
 - if class 1 is encoded to 1 and class 2 is - 1
 - class 1 if $f(x) \geq 0$*
 - class 2 if $f(x) < 0$*
 - linear models are NOT optimized for classification



Logistic regression




Logistic regression

- Predict results on a binary outcome variable
 - e.g., whether or not a patient has a disease
 - whether a new applicant would succeed in the program or not
 - the outcome is not continuous or distributed normally
- when we have a binary response variables
 - we code “disease” as 1 and “no disease” as 0, can we just fit a line through those points as we would with linear regression? **Possible! But some problems.**



Linear Regression Problem

- the problem of fitting a regular regression line to a binary dependent variable
- the line seems to oversimplify the relationship
- it gives predictions that cannot be observable values of Y for extreme values of X
- the approach is analogous to fitting a linear model to the probability of the event
- produces unobservable predictions for extreme values of dependent variable





Probabilistic approach

■ Learn $P(Y|X)$ directly


■ cumulative probability distribution

■ using a sigmoid function $P(Y = 0|X) = \frac{1}{1 + \exp(\mathbf{Xb})}$


$$P(Y = 1|X) = \frac{1}{1 + \exp(-\mathbf{Xb})}$$

$$\text{Sigmoid: } g(w_0 + \sum_i w_i x_i) = \frac{1}{1 + e^{w_0 + \sum_i w_i x_i}}$$


$$w_0 = -2, w_1 = -1$$



$$w_0 = 0, w_1 = -1$$



$$w_0 = 0, w_1 = -0.5$$






Logistic regression model

$$\log \left(\frac{p}{1-p} \right) = \mathbf{Xb}$$


- p is the probability that an event Y occurs, $p(Y=1)$
- $p/(1-p)$ is the "Odds ratio" \rightarrow range of $[0 \text{ to infinite}]$
- $\log(p/(1-p))$ is log odds ratio, or "logit" \rightarrow range: $[-\infty, +\infty]$

Odd	Logit score
0.99	1.996
0.5	0
0.25	-0.477
0.01	-1.996



Logistic function

$$Y = \log(p/(1-p))$$




- Logistic regression predicts probabilities rather than classes:
stochastic approach



Linear Separators

- Binary classification can be viewed as the task of separating classes in feature space:



$$f(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$$

Training set:


$$(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n),$$

y (either 1 or -1) - indicating the class to which the point x_i belongs
 \mathbf{w} - normal vector to the hyperplane



Hyperplane


- lots of possible choices for a , b , c
- a **Support Vector Machine (SVM)** finds an optimal solution
 - maximizes the distance between the hyperplane and the “difficult points” close to decision boundary
 - one intuition: if there are no points near the decision surface, then there are no very uncertain classification decisions





SVM


- SVMs maximize the margin around the separating hyperplane
 - a.k.a. large margin classifiers
- the decision function is fully specified by a subset of training samples, the support vectors
- solving SVMs is a quadratic programming problem





Geometric Margin

- distance from example to the separator is $r = y \frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|}$
- examples closest to the hyperplane are **support vectors**
- margin** ρ of the separator is the width of separation between support vectors of classes



Derivation of finding r :

Dotted line $\mathbf{x}' - \mathbf{x}$ is perpendicular to decision boundary so parallel to \mathbf{w} .

Unit vector is $\mathbf{w}/\|\mathbf{w}\|$, so line is $r\mathbf{w}/\|\mathbf{w}\|$

$\mathbf{x}' = \mathbf{x} - yr\mathbf{w}/\|\mathbf{w}\|$.

\mathbf{x}' satisfies $\mathbf{w}^T \mathbf{x}' + b = 0$.

So $\mathbf{w}^T (\mathbf{x} - yr\mathbf{w}/\|\mathbf{w}\|) + b = 0$

Recall that $\|\mathbf{w}\| = \sqrt{\mathbf{w}^T \mathbf{w}}$.

So $\mathbf{w}^T \mathbf{x} - yr\|\mathbf{w}\| + b = 0$

So, solving for r gives:

$r = y(\mathbf{w}^T \mathbf{x} + b)/\|\mathbf{w}\|$




Linear SVM

- assume that the functional margin of each data item is at least 1, then the following two constraints follow for a training set $\{(x_i, y_i)\}$:

$$\mathbf{w}^T \mathbf{x}_i + b \geq 1 \quad \text{if } y_i = 1$$

$$\mathbf{w}^T \mathbf{x}_i + b \leq -1 \quad \text{if } y_i = -1$$

- for support vectors, the inequality becomes an equality



Hyperplane
 $\mathbf{w}^T \mathbf{x} + b = 0$

Extra scale constraint:
 $\min_{i=1, \dots, n} |\mathbf{w}^T \mathbf{x}_i + b| = 1$

This implies:


$$\mathbf{w}^T (\mathbf{x}_a - \mathbf{x}_b) = 2$$

$$\rho = \|\mathbf{x}_a - \mathbf{x}_b\|_2 = 2 / \|\mathbf{w}\|_2$$



Soft Margin Classification

- if the training data is not linearly separable, slack variables ξ_i can be added to allow misclassification of difficult or noisy examples
- allow some errors
 - let some points be moved to where they belong, at a cost
- still, try to minimize training set errors, and to place hyperplane "far" from each class (large margin)






Classification with SVM

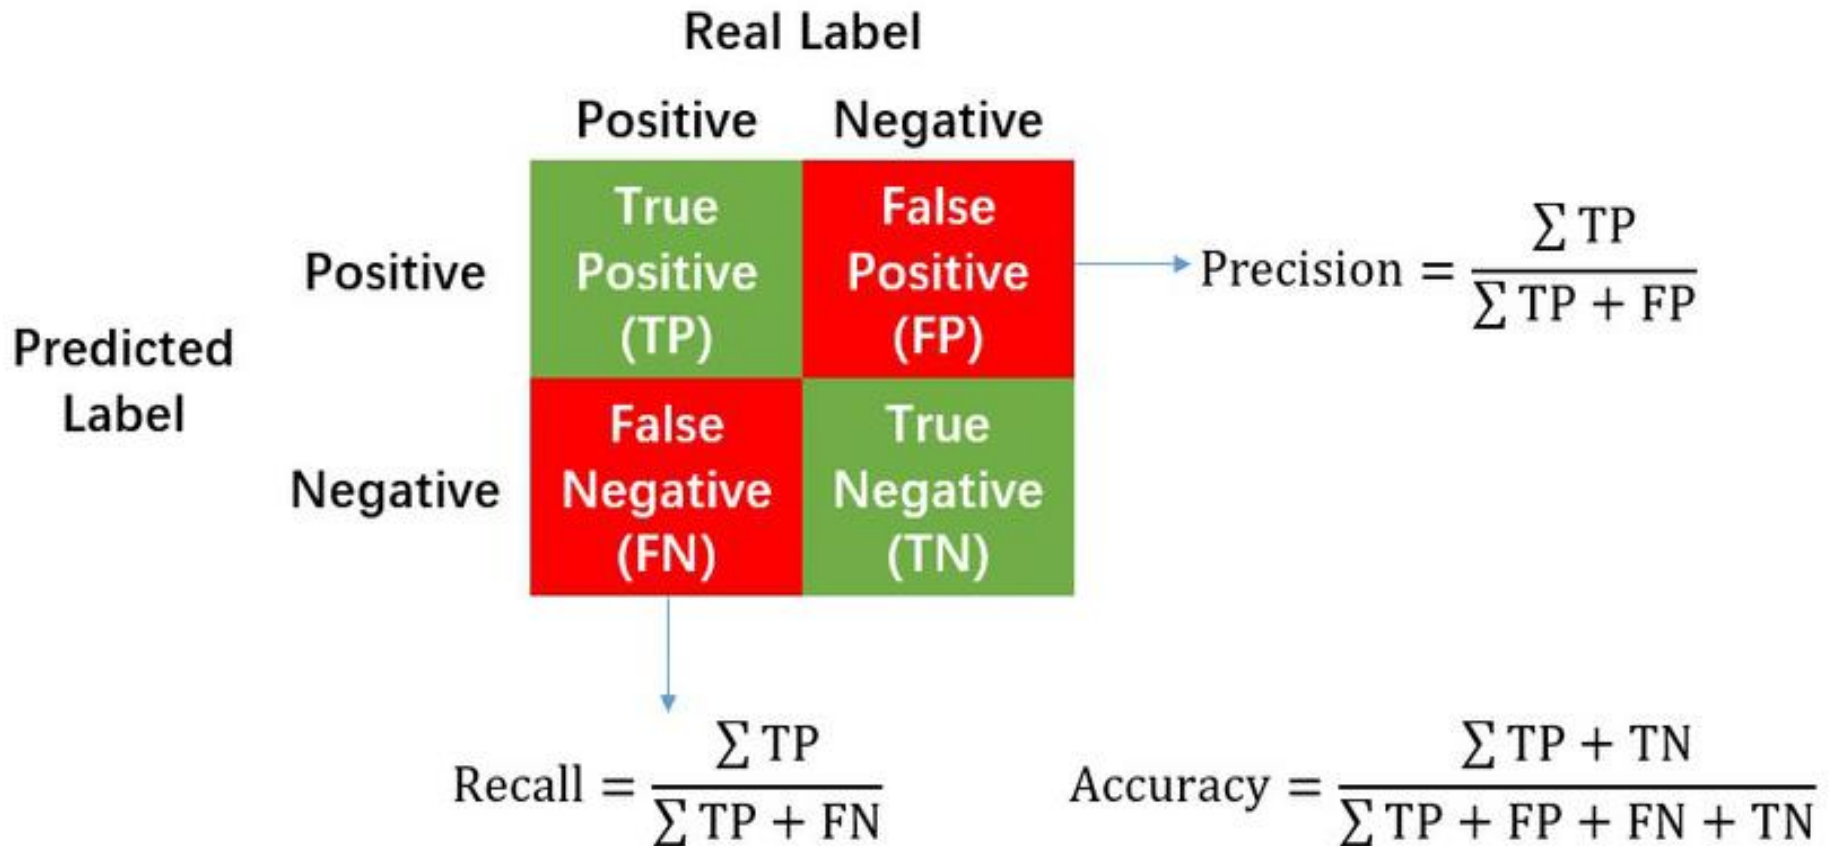
- given a new point x , we can score its projection onto the hyperplane normal:
 - i. e. compute score: $\mathbf{w}^T \mathbf{x} + b = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$
 - decide class based on whether $<$ or $>$ 0
- can set confidence threshold

Score $> t$: **yes**
Score $< -t$: **no**
Else: **don't know**





Metrics





QUESTIONS ?

